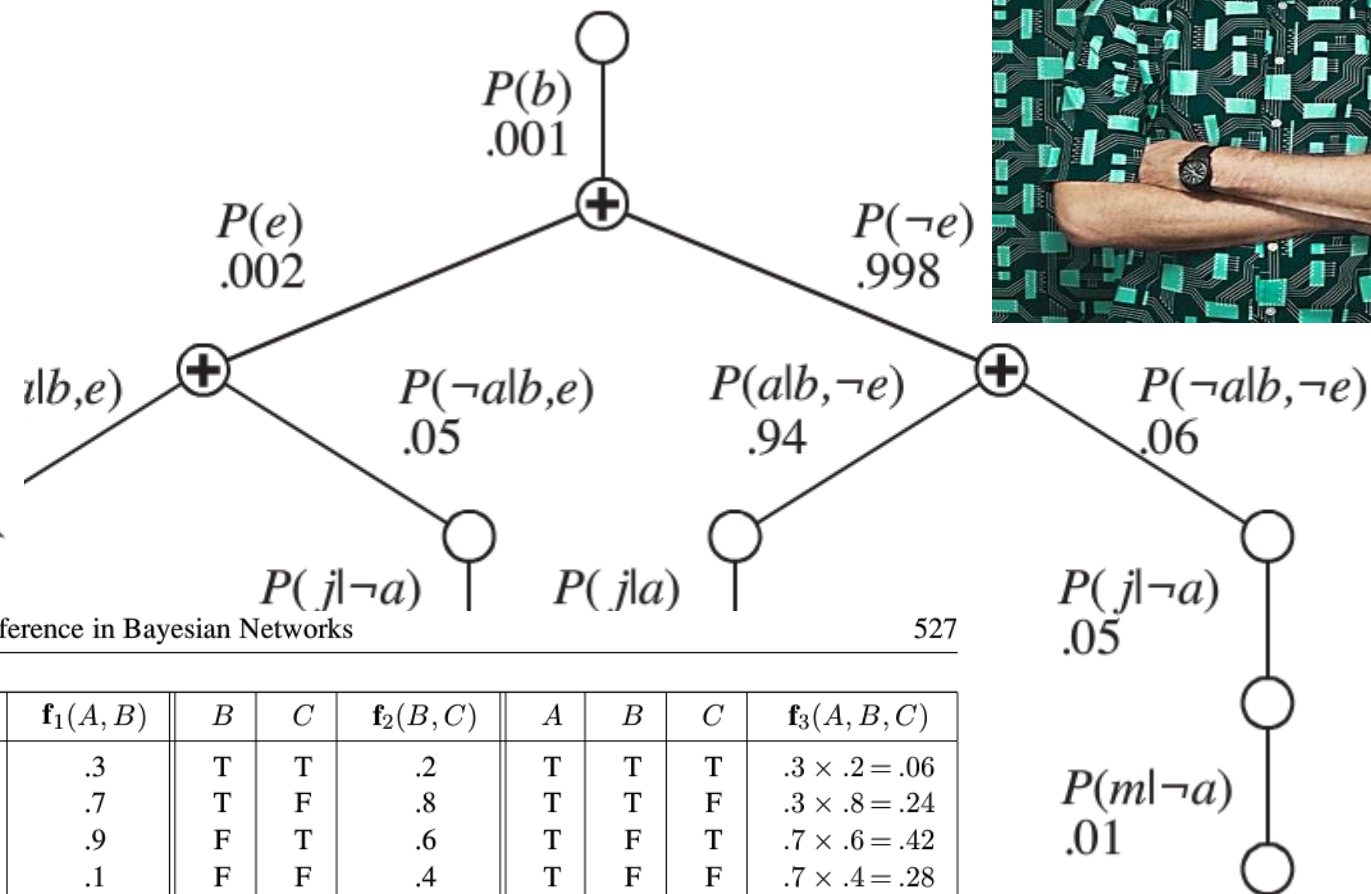






# Principles of AI

## Chapter 14: Probabilistic Reasoning



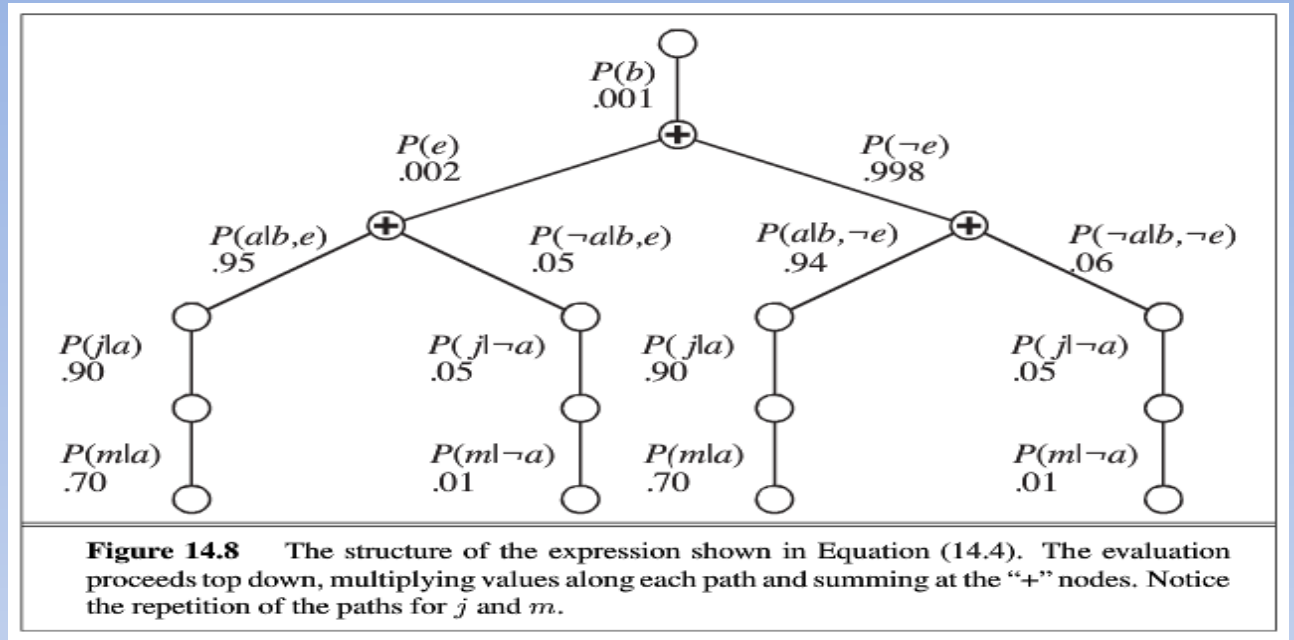
Exact Inference in Bayesian Networks

527

A	B	$f_1(A, B)$	B	C	$f_2(B, C)$	A	B	C	$f_3(A, B, C)$
T	T	.3	T	T	.2	T	T	T	$.3 \times .2 = .06$
T	F	.7	T	F	.8	T	T	F	$.3 \times .8 = .24$
F	T	.9	F	T	.6	T	F	T	$.7 \times .6 = .42$
F	F	.1	F	F	.4	T	F	F	$.7 \times .4 = .28$
						F	T	T	$.9 \times .2 = .18$
						F	T	F	$.9 \times .8 = .72$
						F	F	T	$.1 \times .6 = .06$
						F	F	F	$.1 \times .4 = .04$

**Figure 14.10** Illustrating pointwise multiplication:  $f_1(A, B) \times f_2(B, C) = f_3(A, B, C)$ .

# Exact Inference: Enumeration



- Recursive depth-first enumeration:
  - $O(n)$  space,
  - $O(d^n)$  time
- Lots of repeated calculations
- Maybe Dynamic Programming!

# Variable Elimination

- Variable elimination:
  - carry out summations right-to-left,
  - store intermediate results (factors) to avoid recomputation

**function** ELIMINATION-ASK( $X, \mathbf{e}, bn$ ) **returns** a distribution over  $X$   
**inputs:**  $X$ , the query variable  
 $\mathbf{e}$ , observed values for variables  $\mathbf{E}$   
 $bn$ , a Bayesian network specifying joint distribution  $\mathbf{P}(X_1, \dots, X_n)$

$factors \leftarrow []$   
**for each**  $var$  **in** ORDER( $bn.VARS$ ) **do**  
 $factors \leftarrow [MAKE-FACTOR(var, \mathbf{e}) | factors]$   
**if**  $var$  is a hidden variable **then**  $factors \leftarrow SUM-OUT(var, factors)$   
**return** NORMALIZE(POINTWISE-PRODUCT( $factors$ ))

**Procedure** Sum-Product-VE (

$\Phi$ , // Set of factors  
 $Z$ , // Set of variables to be eliminated  
 $\prec$  // Ordering on  $Z$

)

1 Let  $Z_1, \dots, Z_k$  be an ordering of  $Z$  such that  
2  $Z_i \prec Z_j$  if and only if  $i < j$   
3 **for**  $i = 1, \dots, k$   
4  $\Phi \leftarrow \text{Sum-Product-Eliminate-Var}(\Phi, Z_i)$   
5  $\phi^* \leftarrow \prod_{\phi \in \Phi} \phi$   
6 **return**  $\phi^*$

**Procedure** Sum-Product-Eliminate-Var (

$\Phi$ , // Set of factors  
 $Z$  // Variable to be eliminated

)

1  $\Phi' \leftarrow \{\phi \in \Phi : Z \in \text{Scope}[\phi]\}$   
2  $\Phi'' \leftarrow \Phi - \Phi'$   
3  $\psi \leftarrow \prod_{\phi \in \Phi'} \phi$   
4  $\tau \leftarrow \sum_Z \psi$   
5 **return**  $\Phi'' \cup \{\tau\}$

**function** ELIMINATION-ASK( $X, \mathbf{e}, bn$ ) **returns** a distribution over  $X$   
**inputs:**  $X$ , the query variable  
 $\mathbf{e}$ , observed values for variables  $\mathbf{E}$   
 $bn$ , a Bayesian network specifying joint distribution  $\mathbf{P}(X_1, \dots, X_n)$

304

Chapter 9. Variable Elimination

**Procedure** Sum-Product-VE (

$\Phi$ , // Set of factors

$Z$ , // Set of variables to be eliminated

$\prec$  // Ordering on  $Z$

)

1 Let  $Z_1, \dots, Z_k$  be an ordering of  $Z$

2  $Z_i \prec Z_j$  if and only if  $i < j$

3 **for**  $i = 1, \dots, k$

4  $\Phi \leftarrow \text{Sum-Product-Eliminate}(\Phi, Z_i, \prec)$

5  $\phi^* \leftarrow \prod_{\phi \in \Phi} \phi$

6 **return**  $\phi^*$

**Procedure** Sum-Product-Eliminate (

$\Phi$ , // Set of factors

$Z$  // Variable to be eliminated

)

1  $\Phi' \leftarrow \{\phi \in \Phi : Z \in \text{Scope}[\phi]\}$

2  $\Phi'' \leftarrow \Phi - \Phi'$

3  $\psi \leftarrow \prod_{\phi \in \Phi'} \phi$

4  $\tau \leftarrow \sum_Z \psi$

5 **return**  $\Phi'' \cup \{\tau\}$

**Algorithm 9.2 Using Sum-Product-VE for computing conditional probabilities**

**Procedure** Cond-Prob-VE (

$\mathcal{K}$ , // A network over  $\mathcal{X}$

$\mathbf{Y}$ , // Set of query variables

$\mathbf{E} = \mathbf{e}$  // Evidence

)

1  $\Phi \leftarrow$  Factors parameterizing  $\mathcal{K}$

2 Replace each  $\phi \in \Phi$  by  $\phi[\mathbf{E} = \mathbf{e}]$

3 Select an elimination ordering  $\prec$

4  $Z \leftarrow \mathcal{X} - \mathbf{Y} - \mathbf{E}$

5  $\phi^* \leftarrow \text{Sum-Product-VE}(\Phi, \prec, Z)$

6  $\alpha \leftarrow \sum_{\mathbf{y} \in \text{Val}(\mathbf{Y})} \phi^*(\mathbf{y})$

7 **return**  $\alpha, \phi^*$




# Eliminate One Hidden Variable

**Procedure** Sum-Product-Eliminate-Var (  
     $\Phi$ ,   // Set of factors  
     $Z$     // Variable to be eliminated  
)


```
1    $\Phi' \leftarrow \{\phi \in \Phi : Z \in \text{Scope}[\phi]\}$   
2    $\Phi'' \leftarrow \Phi - \Phi'$   
3    $\psi \leftarrow \prod_{\phi \in \Phi'} \phi$   
4    $\tau \leftarrow \sum_Z \psi$   
5   return  $\Phi'' \cup \{\tau\}$ 
```

## 2 Variable Elimination

For the Bayes' net below, we are given the query  $P(Y \mid +z)$ . All variables have binary domains. Assume we run variable elimination to compute the answer to this query, with the following variable elimination ordering:  $X, T, U, V, W$ .



# $P(Y \mid +z)$ : Eliminate Hidden Variables




- Initial Factors after inserting evidence:
- $P(T)$ ,  $P(U|T)$ ,  $P(V|T)$ ,  $P(W|T)$ ,  $P(X|T)$ ,  $P(Y|V,W)$ ,  $P(+z|X)$



$$P(Y \mid +z)$$


Eliminate First Hidden Variable:  
**X**



- $P(T), P(U|T), P(V|T), P(W|T), P(X|T), P(Y|V,W), P(+z|X)$
- (a) Now Eliminate X and generate a new factor f1:

$$P(Y \mid +z)$$

- $P(T)$ ,  $P(U|T)$ ,  $P(V|T)$ ,  $P(W|T)$ ,  $P(X|T)$ ,  $P(Y|V,W)$ ,  $P(+z|X)$
- (a) Now Eliminate  $X$  and generate a new factor  $f_1$ :
  - Step 1: Collect/Combine Factors
    - $P(X|T)$ ,  $P(+z|X)$
  - Step 2: Marginalize out  $X$ 
    - $f_1(T, +z) = \sum_x P(x|T)P(+z|x)$
  - (b) Leaving us w/ Factors =  $P(T)$ ,  $P(U|T)$ ,  $P(V|T)$ ,  $P(W|T)$ ,  $P(Y|V,W)$ ,  **$f_1(T, +z)$**




$$P(Y \mid +z)$$

Eliminate Second Hidden Variable:

**T**


- $P(T), P(U|T), P(V|T), P(W|T), P(Y|V,W), f_1(T,+z)$
- (c) Eliminate T -- generating a new factor f2:
  - Step 1: Collect/Combine Factors
    - $P(T), P(U|T), P(V|T), P(W|T), f_1(T,+z)$
  - Step 2: Marginalize out T
    - $f_2(U, V, W, +z)$   
 $= \sum_t P(T), P(U|T), P(V|T), P(W|T), f_1(T,+z)$
  - (d) Leaving us w/ Factors=  $P(Y|V,W), f_2(U,V,W,+z)$



$$P(Y \mid +z)$$

## Eliminate Third Hidden Variable: **U**


- $P(Y \mid V, W), f_2(U, V, W, +z)$
- (e) Eliminate U, generate a new factor f3 :
  - Step 1: Collect/Combine Factors
    - $f_2(U, V, W, +z)$
  - Step 2: Marginalize out T
    - $f_3(V, W, +z) = \sum_u f_2(u, V, W, +z)$
  - (f) Leaving us w/ Factors=  $P(Y \mid V, W), f_3(V, W, +z)$



$$P(Y \mid +z)$$

## Eliminate Fourth Hidden Variable: **V**


- $P(Y \mid V, W), f_3(V, W, +z)$
- (g) Eliminating  $V$ , generate a new factor  $f_4$ :
  - Step 1: Collect/Combine Factors
    - $P(Y \mid V, W), f_3(V, W, +z)$
  - Step 2: Marginalize out  $V$ 
    - $f_4(Y, W, +z) = \sum_v P(Y \mid v, W), f_3(v, W, +z)$
  - (h) Leaving us w/ Factors =  $f_4(Y, W, +z)$



$$P(Y \mid +z)$$


Eliminate Fifth Hidden Variable:  
**W**

- $f_4(Y, W, +z)$
- (i) Eliminating  $W$ , generate a new factor  $f_5$ :
  - Step 1: Collect/Combine Factors
    - $f_4(Y, W, +z)$
  - Step 2: Marginalize out  $W$ 
    - $f_5(Y, +z) = \sum_w f_4(Y, W, +z)$
  - (j) Leaving us w/ Factors =  $f_5(Y, +z)$




$$P(Y \mid +z)$$

- $f_5(Y, +z)$
- (k) How would you obtain  $P(Y \mid +z)$  from the factors left above:
  - Simply renormalize  $f_5(Y, +z)$  to obtain  $P(Y \mid +z)$ :
    - $$P(Y \mid +z) = \frac{f_5(y, +z)}{\sum_{y'} f_5(y', +z)}$$
- (l) What is the size of the largest factor generated?
- (m) Does there exist a better elimination ordering (one which generates smaller largest factors)?




$$P(Y \mid +z)$$

- $f_5(Y, +z)$
- (k) How would you obtain  $P(Y \mid +z)$  from the factors left above:
  - Simply renormalize  $f_5(Y, +z)$  to obtain  $P(Y \mid +z)$ :
    - $$P(Y \mid +z) = \frac{f_5(y, +z)}{\sum_{y'} f_5(y', +z)}$$
- (l) What is the size of the largest factor generated?
- (m) Does there exist a better elimination ordering (one which generates smaller largest factors)? **[X, U, V, T, W]**






# Approximate Inference: Sampling



# Sampling

- Sampling is a lot like repeated simulation
    - Predicting the weather, basketball games, ...
  - Basic idea
    - Draw  $N$  samples from a sampling distribution  $S$
    - Compute an approximate posterior probability
    - Show this converges to the true probability  $P$
- Why sample?
    - Learning: get samples from a distribution you don't know
    - Inference: getting a sample is faster than computing the right answer (e.g. with variable elimination)



# Sampling


- Sampling from given distribution
  - Step 1: Get sample  $u$  from uniform distribution over  $[0, 1)$ 
    - E.g. `random()` in python
  - Step 2: Convert this sample  $u$  into an outcome for the given distribution by having each outcome associated with a sub-interval of  $[0,1)$  with sub-interval size equal to probability of the outcome

## ■ Example

C	P(C)
red	0.6
green	0.1
blue	0.3

$0 \leq u < 0.6, \rightarrow C = \text{red}$   
 $0.6 \leq u < 0.7, \rightarrow C = \text{green}$   
 $0.7 \leq u < 1, \rightarrow C = \text{blue}$


- If `random()` returns  $u = 0.83$ , then our sample is  $C = \text{blue}$
- E.g, after sampling 8 times:



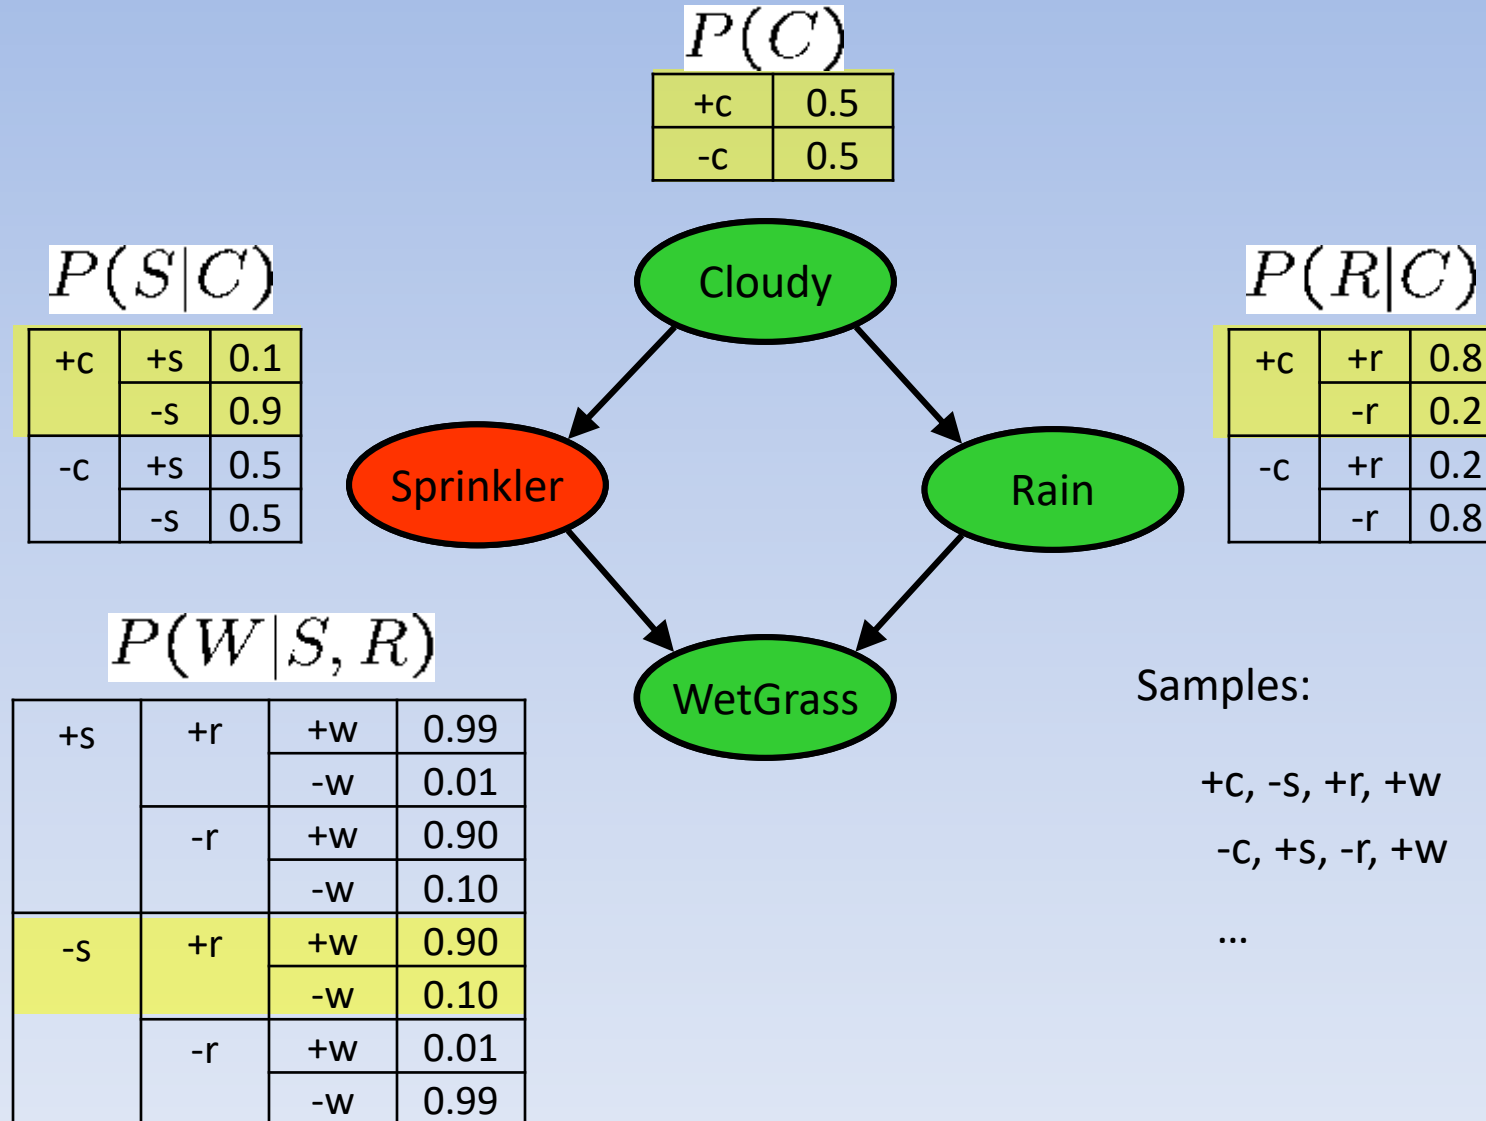
# Sampling in Bayes' Nets

- Prior Sampling
- Rejection Sampling
- Likelihood Weighting
- Gibbs Sampling

# Prior Sampling




# Prior Sampling



# Prior Sampling

- For  $i=1, 2, \dots, n$ 
  - Sample  $x_i$  from  $P(X_i \mid \text{Parents}(X_i))$
- Return  $(x_1, x_2, \dots, x_n)$



```
function PRIOR-SAMPLE(bn) returns an event sampled from the prior specified by bn  
  inputs: bn, a Bayesian network specifying joint distribution  $\mathbf{P}(X_1, \dots, X_n)$   
  
   $\mathbf{x} \leftarrow$  an event with  $n$  elements  
  foreach variable  $X_i$  in  $X_1, \dots, X_n$  do  
     $\mathbf{x}[i] \leftarrow$  a random sample from  $\mathbf{P}(X_i \mid \text{parents}(X_i))$   
  return  $\mathbf{x}$ 
```

**Figure 14.13** A sampling algorithm that generates events from a Bayesian network. Each variable is sampled according to the conditional distribution given the values already sampled for the variable's parents.



## Sampling from an empty network

```
function PRIOR-SAMPLE(bn) returns an event sampled from bn
  inputs: bn, a belief network specifying joint distribution  $P(X_1, \dots, X_n)$ 
   $x \leftarrow$  an event with  $n$  elements
  for  $i = 1$  to  $n$  do
     $x_i \leftarrow$  a random sample from  $P(X_i \mid \text{parents}(X_i))$ 
      given the values of  $\text{Parents}(X_i)$  in  $x$ 
  return  $x$ 
```

# Prior Sampling

- This process generates samples with probability:

$$S_{PS}(x_1 \dots x_n) = \prod_{i=1}^n P(x_i | \text{Parents}(X_i)) = P(x_1 \dots x_n)$$

...i.e. the BN's joint probability

- Let the number of samples of an event be  $N_{PS}(x_1 \dots x_n)$
- Then 
$$\begin{aligned} \lim_{N \rightarrow \infty} \hat{P}(x_1, \dots, x_n) &= \lim_{N \rightarrow \infty} N_{PS}(x_1, \dots, x_n) / N \\ &= S_{PS}(x_1, \dots, x_n) \\ &= P(x_1 \dots x_n) \end{aligned}$$
- I.e., the sampling procedure is **consistent**

# Example

- We'll get a bunch of samples from the BN:

+c, -s, +r, +w

+c, +s, +r, +w


-c, +s, +r, -w

+c, -s, +r, +w


-c, -s, -r, +w

- If we want to know  $P(W)$

- We have counts  $\langle +w:4, -w:1 \rangle$
- Normalize to get  $P(W) = \langle +w:0.8, -w:0.2 \rangle$
- This will get closer to the true distribution with more samples
- Can estimate anything else, too
- What about  $P(C \mid +w)$ ?  $P(C \mid +r, +w)$ ?  $P(C \mid -r, -w)$ ?
- Fast: can use fewer samples if less time (what's the drawback?)




# Rejection Sampling



# Rejection Sampling


- Let's say we want  $P(C)$ 
  - No point keeping all samples around
  - Just tally counts of  $C$  as we go
- Let's say we want  $P(C \mid +s)$ 
  - Same thing: tally  $C$  outcomes, but ignore (reject) samples which don't have  $S=+s$
  - This is called rejection sampling
  - It is also consistent for conditional probabilities (i.e., correct in the limit)



+c, -s, +r, +w  
+c, +s, +r, +w  
-c, +s, +r, -w  
+c, -s, +r, +w  
-c, -s, -r, +w

# Rejection Sampling

- IN: evidence instantiation
- For  $i=1, 2, \dots, n$ 
  - Sample  $x_i$  from  $P(X_i \mid \text{Parents}(X_i))$
  - If  $x_i$  not consistent with evidence
    - Reject: Return, and no sample is generated in this cycle
- Return  $(x_1, x_2, \dots, x_n)$




```
function REJECTION-SAMPLING( $X, \mathbf{e}, bn, N$ ) returns an estimate of  $\mathbf{P}(X|\mathbf{e})$   
  inputs:  $X$ , the query variable  
            $\mathbf{e}$ , observed values for variables  $\mathbf{E}$   
            $bn$ , a Bayesian network  
            $N$ , the total number of samples to be generated  
  local variables:  $\mathbf{N}$ , a vector of counts for each value of  $X$ , initially zero  
  
  for  $j = 1$  to  $N$  do  
     $\mathbf{x} \leftarrow$  PRIOR-SAMPLE( $bn$ )  
    if  $\mathbf{x}$  is consistent with  $\mathbf{e}$  then  
       $\mathbf{N}[x] \leftarrow \mathbf{N}[x] + 1$  where  $x$  is the value of  $X$  in  $\mathbf{x}$   
  return NORMALIZE( $\mathbf{N}$ )
```

**Figure 14.14** The rejection-sampling algorithm for answering queries given evidence in a Bayesian network.




# Likelihood Weighting





# Likelihood Weighting

- Problem with rejection sampling:
  - If evidence is unlikely, rejects lots of samples
  - Evidence not exploited as you sample
  - Consider  $P(\text{Shape} | \text{blue})$




~~pyramid, green~~  
~~pyramid, red~~  
sphere, blue  
~~cube, red~~  
~~sphere, green~~



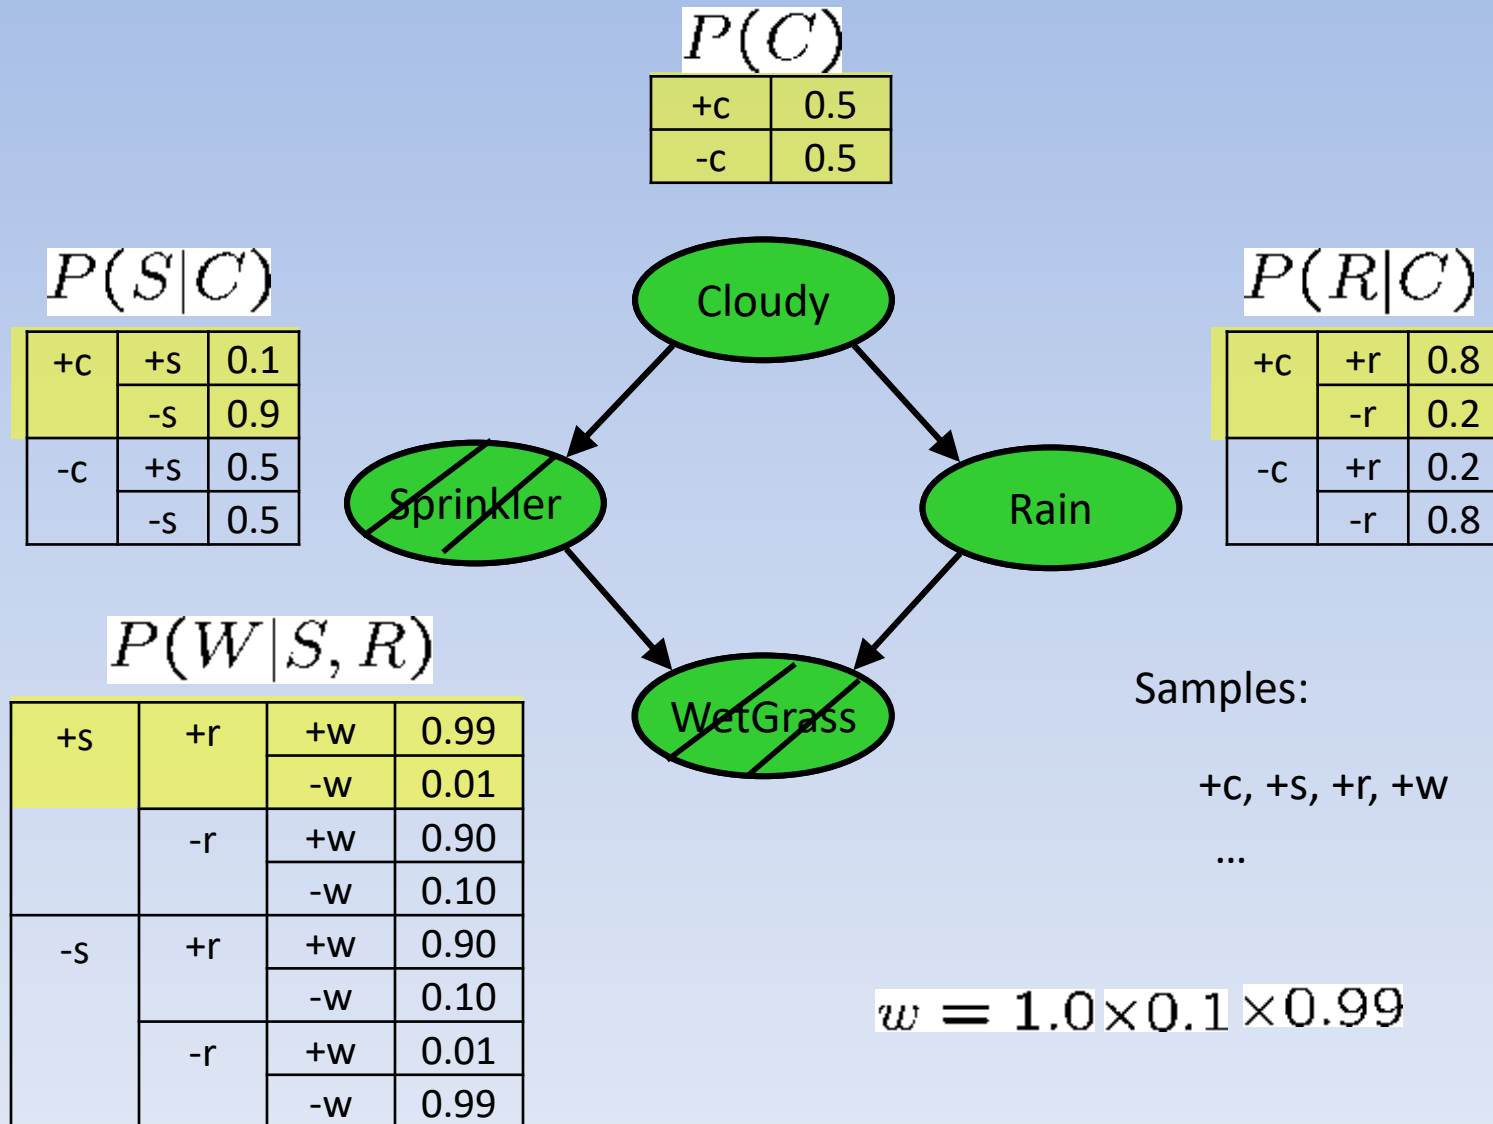
- Idea: fix evidence variables and sample the rest
  - Problem: sample distribution not consistent!
  - Solution: weight by probability of evidence given parents



pyramid, blue  
pyramid, blue  
sphere, blue  
cube, blue  
sphere, blue




# Likelihood Weighting



# Likelihood Weighting

- IN: evidence instantiation
- $w = 1.0$
- for  $i=1, 2, \dots, n$ 
  - if  $X_i$  is an evidence variable
    - $X_i = \text{observation } x_i \text{ for } X_i$
    - Set  $w = w * P(x_i \mid \text{Parents}(X_i))$
  - else
    - Sample  $x_i$  from  $P(X_i \mid \text{Parents}(X_i))$
- return  $(x_1, x_2, \dots, x_n), w$



```

function LIKELIHOOD-WEIGHTING( $X, \mathbf{e}, bn, N$ ) returns an estimate of  $\mathbf{P}(X|\mathbf{e})$ 
  inputs:  $X$ , the query variable
            $\mathbf{e}$ , observed values for variables  $\mathbf{E}$ 
            $bn$ , a Bayesian network specifying joint distribution  $\mathbf{P}(X_1, \dots, X_n)$ 
            $N$ , the total number of samples to be generated
  local variables:  $\mathbf{W}$ , a vector of weighted counts for each value of  $X$ , initially zero

  for  $j = 1$  to  $N$  do
     $\mathbf{x}, w \leftarrow \text{WEIGHTED-SAMPLE}(bn, \mathbf{e})$ 
     $\mathbf{W}[x] \leftarrow \mathbf{W}[x] + w$  where  $x$  is the value of  $X$  in  $\mathbf{x}$ 
  return NORMALIZE( $\mathbf{W}$ )

```

---

```

function WEIGHTED-SAMPLE( $bn, \mathbf{e}$ ) returns an event and a weight

   $w \leftarrow 1$ ;  $\mathbf{x} \leftarrow$  an event with  $n$  elements initialized from  $\mathbf{e}$ 
  foreach variable  $X_i$  in  $X_1, \dots, X_n$  do
    if  $X_i$  is an evidence variable with value  $x_i$  in  $\mathbf{e}$ 
      then  $w \leftarrow w \times P(X_i = x_i \mid \text{parents}(X_i))$ 
      else  $\mathbf{x}[i] \leftarrow$  a random sample from  $\mathbf{P}(X_i \mid \text{parents}(X_i))$ 
  return  $\mathbf{x}, w$ 

```

**Figure 14.15** The likelihood-weighting algorithm for inference in Bayesian networks. In WEIGHTED-SAMPLE, each nonevidence variable is sampled according to the conditional distribution given the values already sampled for the variable's parents, while a weight is accumulated based on the likelihood for each evidence variable.


# Likelihood Weighting

- Sampling distribution if  $z$  sampled and  $e$  fixed evidence

$$S_{WS}(z, e) = \prod_{i=1}^l P(z_i | \text{Parents}(Z_i))$$

- Now, samples have weights

$$w(z, e) = \prod_{i=1}^m P(e_i | \text{Parents}(E_i))$$




- Together, weighted sampling distribution is consistent


$$\begin{aligned} S_{WS}(z, e) \cdot w(z, e) &= \prod_{i=1}^l P(z_i | \text{Parents}(z_i)) \prod_{i=1}^m P(e_i | \text{Parents}(e_i)) \\ &= P(z, e) \end{aligned}$$

# Likelihood Weighting

- Likelihood weighting is good
  - We have taken evidence into account as we generate the sample
  - E.g. here,  $W$ 's value will get picked based on the evidence values of  $S$ ,  $R$
  - More of our samples will reflect the state of the world suggested by the evidence
- Likelihood weighting doesn't solve all our problems
  - Evidence influences the choice of downstream variables, but not upstream ones ( $C$  isn't more likely to get a value matching the evidence)
- We would like to consider evidence when we sample every variable
  - Gibbs sampling



# Gibbs Sampling





# Gibbs Sampling


- *Procedure:* keep track of a full instantiation  $x_1, x_2, \dots, x_n$ .
  - Start with an arbitrary instantiation consistent with the evidence.
  - Sample one variable at a time, conditioned on all the rest, but keep evidence fixed.
  - Keep repeating this for a long time.
- *Property:* in the limit of repeating this infinitely many times the resulting sample is coming from the correct distribution
- *Rationale:* both upstream and downstream variables condition on evidence.
- **In contrast:**
  - likelihood weighting only conditions on upstream evidence,
  - hence weights obtained in likelihood weighting can sometimes be very small.
  - Sum of weights over all samples is indicative of how many “effective” samples were obtained,
    - so want high weight.



# Gibbs Sampling Example: $P(S \mid +r)$


- Step 1: Fix evidence

- $R = +r$




- Step 2: Initialize other variables

- Randomly



- Steps 3: Repeat

- Choose a non-evidence variable  $X$
  - Resample  $X$  from  $P(X \mid \text{all other variables})$



Sample from  $P(S \mid +c, -w, +r)$

Sample from  $P(C \mid +s, -w, +r)$

Sample from  $P(W \mid +s, +c, +r)$

```
function GIBBS-ASK( $X, \mathbf{e}, bn, N$ ) returns an estimate of  $\mathbf{P}(X|\mathbf{e})$   
  local variables:  $\mathbf{N}$ , a vector of counts for each value of  $X$ , initially zero  
                    $\mathbf{Z}$ , the nonevidence variables in  $bn$   
                    $\mathbf{x}$ , the current state of the network, initially copied from  $\mathbf{e}$   
  
  initialize  $\mathbf{x}$  with random values for the variables in  $\mathbf{Z}$   
  for  $j = 1$  to  $N$  do  
    for each  $Z_i$  in  $\mathbf{Z}$  do  
      set the value of  $Z_i$  in  $\mathbf{x}$  by sampling from  $\mathbf{P}(Z_i|mb(Z_i))$   
       $\mathbf{N}[x] \leftarrow \mathbf{N}[x] + 1$  where  $x$  is the value of  $X$  in  $\mathbf{x}$   
  return NORMALIZE( $\mathbf{N}$ )
```

**Figure 14.16** The Gibbs sampling algorithm for approximate inference in Bayesian networks; this version cycles through the variables, but choosing variables at random also works.

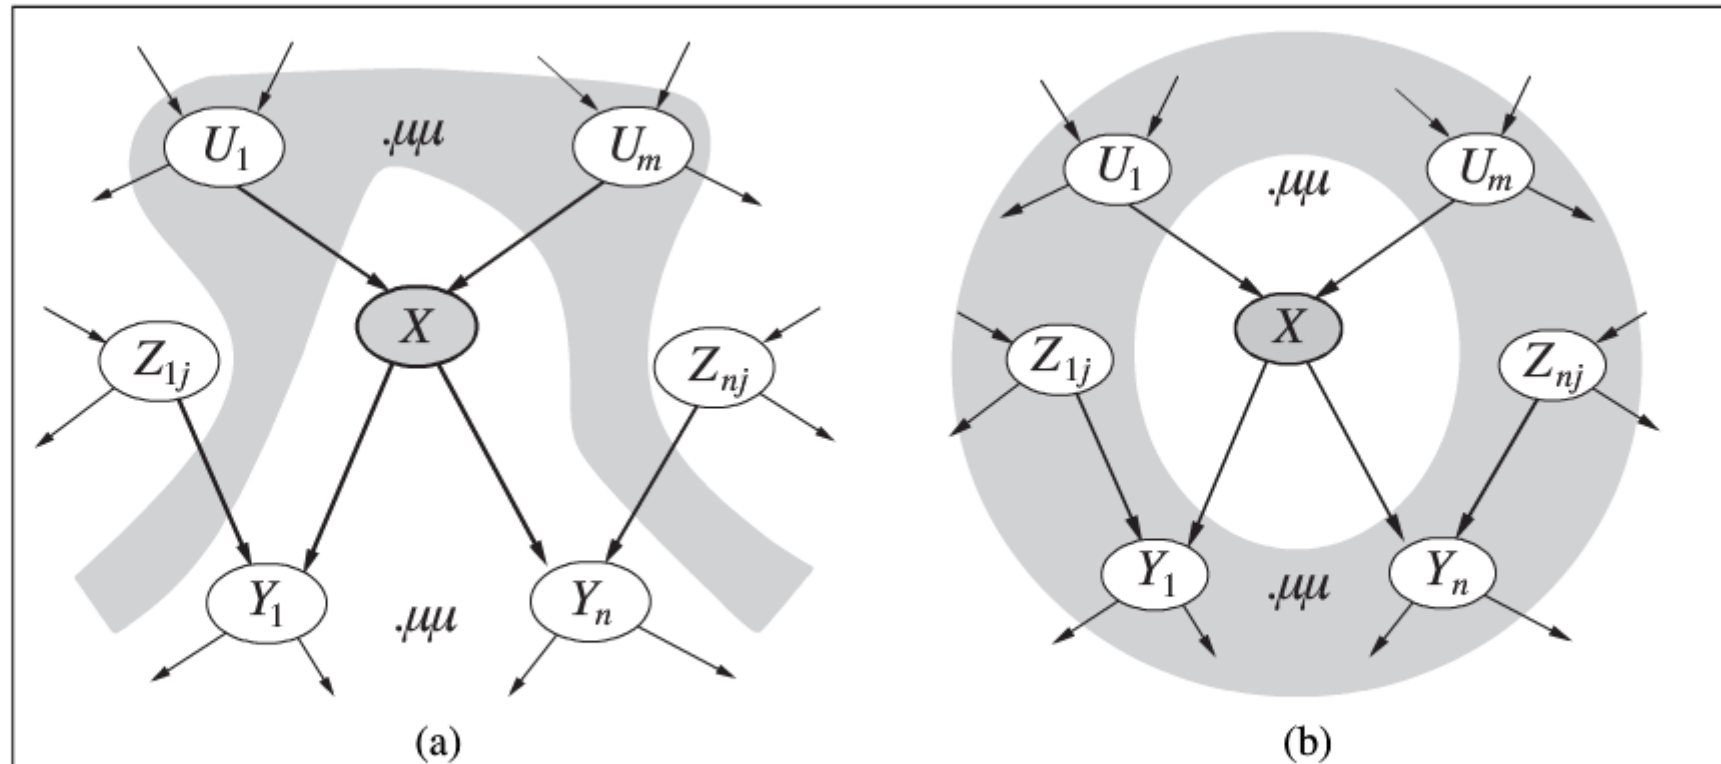
# Gibbs Sampling

- How is this better than sampling from the full joint?
  - In a Bayes' Net, sampling a variable given all the other variables (e.g.  $P(R|S,C,W)$ ) is usually much easier than sampling from the full joint distribution
    - Only requires a join on the variable to be sampled (in this case, a join on R)
    - The resulting factor only depends on the variable's parents, its children, and its children's parents (this is often referred to as its Markov blanket)

# Markov Blanket

518

Chapter 14. Probabilistic Reasoning




**Figure 14.4** (a) A node  $X$  is conditionally independent of its non-descendants (e.g., the  $Z_{ij}$ s) given its parents (the  $U_i$ s shown in the gray area). (b) A node  $X$  is conditionally independent of all other nodes in the network given its Markov blanket (the gray area).

# Efficient Resampling of One Variable

- Sample from  $P(S \mid +c, +r, -w)$


$$\begin{aligned} P(S \mid +c, +r, -w) &= \frac{P(S, +c, +r, -w)}{P(+c, +r, -w)} \\ &= \frac{P(S, +c, +r, -w)}{\sum_s P(s, +c, +r, -w)} \\ &= \frac{P(+c)P(S \mid +c)P(+r \mid +c)P(-w \mid S, +r)}{\sum_s P(+c)P(s \mid +c)P(+r \mid +c)P(-w \mid s, +r)} \\ &= \frac{P(+c)P(S \mid +c)P(+r \mid +c)P(-w \mid S, +r)}{P(+c)P(+r \mid +c) \sum_s P(s \mid +c)P(-w \mid s, +r)} \\ &= \frac{P(S \mid +c)P(-w \mid S, +r)}{\sum_s P(s \mid +c)P(-w \mid s, +r)} \end{aligned}$$

- Many things cancel out – only CPTs with  $S$  remain!
- More generally: only CPTs that have resampled variable need to be considered, and joined together



# Bayes' Net Sampling Summary


- Prior Sampling  $P$



- Rejection Sampling  $P(Q | e)$



- Likelihood Weighting  $P(Q | e)$



- Gibbs Sampling  $P(Q | e)$




# Further Reading on Gibbs Sampling\*

- Gibbs sampling produces sample from the query distribution  $P(Q | e)$  in limit of re-sampling infinitely often
- Gibbs sampling is a special case of more general methods called Markov chain Monte Carlo (MCMC) methods
  - Metropolis-Hastings is one of the more famous MCMC methods (in fact, Gibbs sampling is a special case of Metropolis-Hastings)
- You may read about Monte Carlo methods – they're just sampling

# Markov Chain Monte Carlo\*

- *Idea*: instead of sampling from scratch, create samples that are each like the last one.
- *Procedure*: resample one variable at a time, conditioned on all the rest, but keep evidence fixed. E.g., for  $P(b | c)$ :



- *Properties*: Now samples are not independent (in fact they're nearly identical), but sample averages are still consistent estimators!
- *What's the point*: both upstream and downstream variables condition on evidence.